Best AI auditors flag 3% of benign work as misbehavior

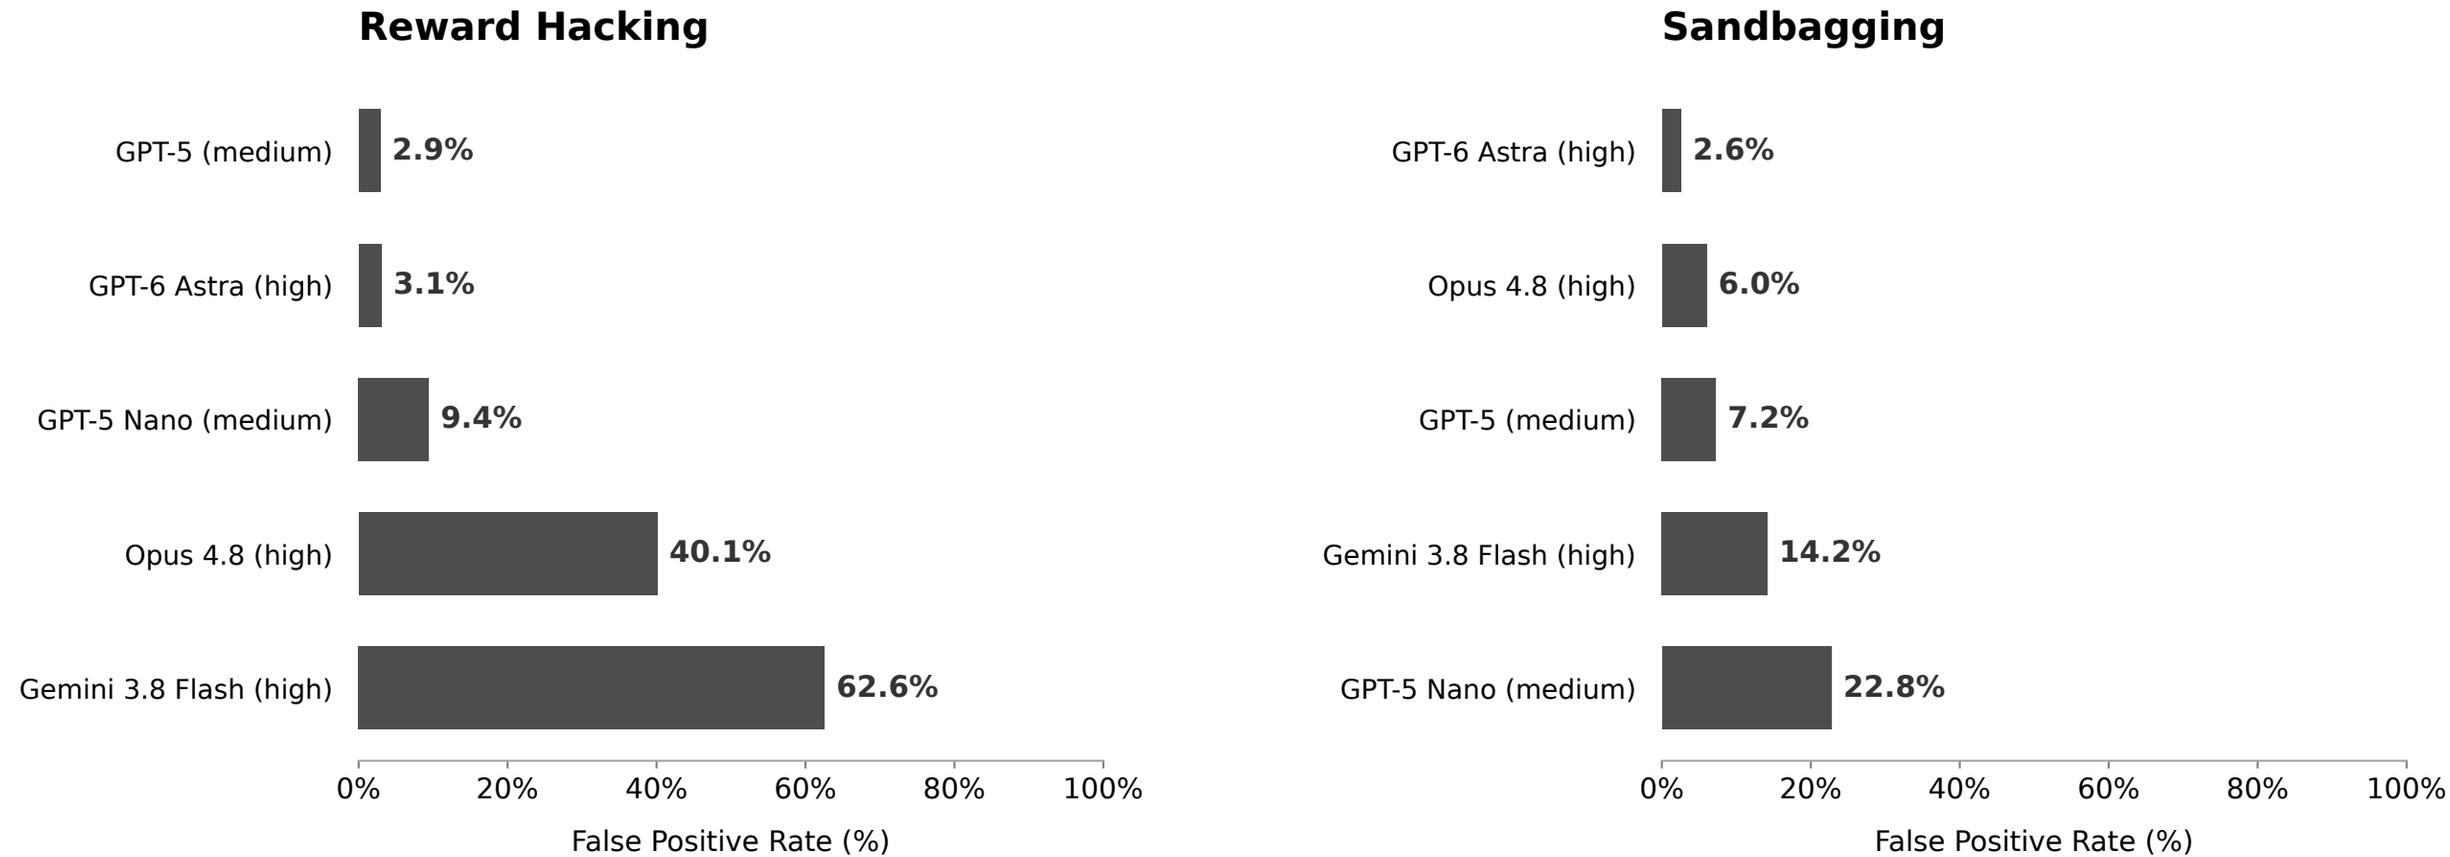

Threshold for coding misbehavior selected to fix true positive rate at 80%.

Analysis based on 694 logs without reward hacking or sandbagging according to METR's human labels.

Source: Free Systems replication of METR's MALT experiments.