AI auditors find over 80% of cases of misbehavior

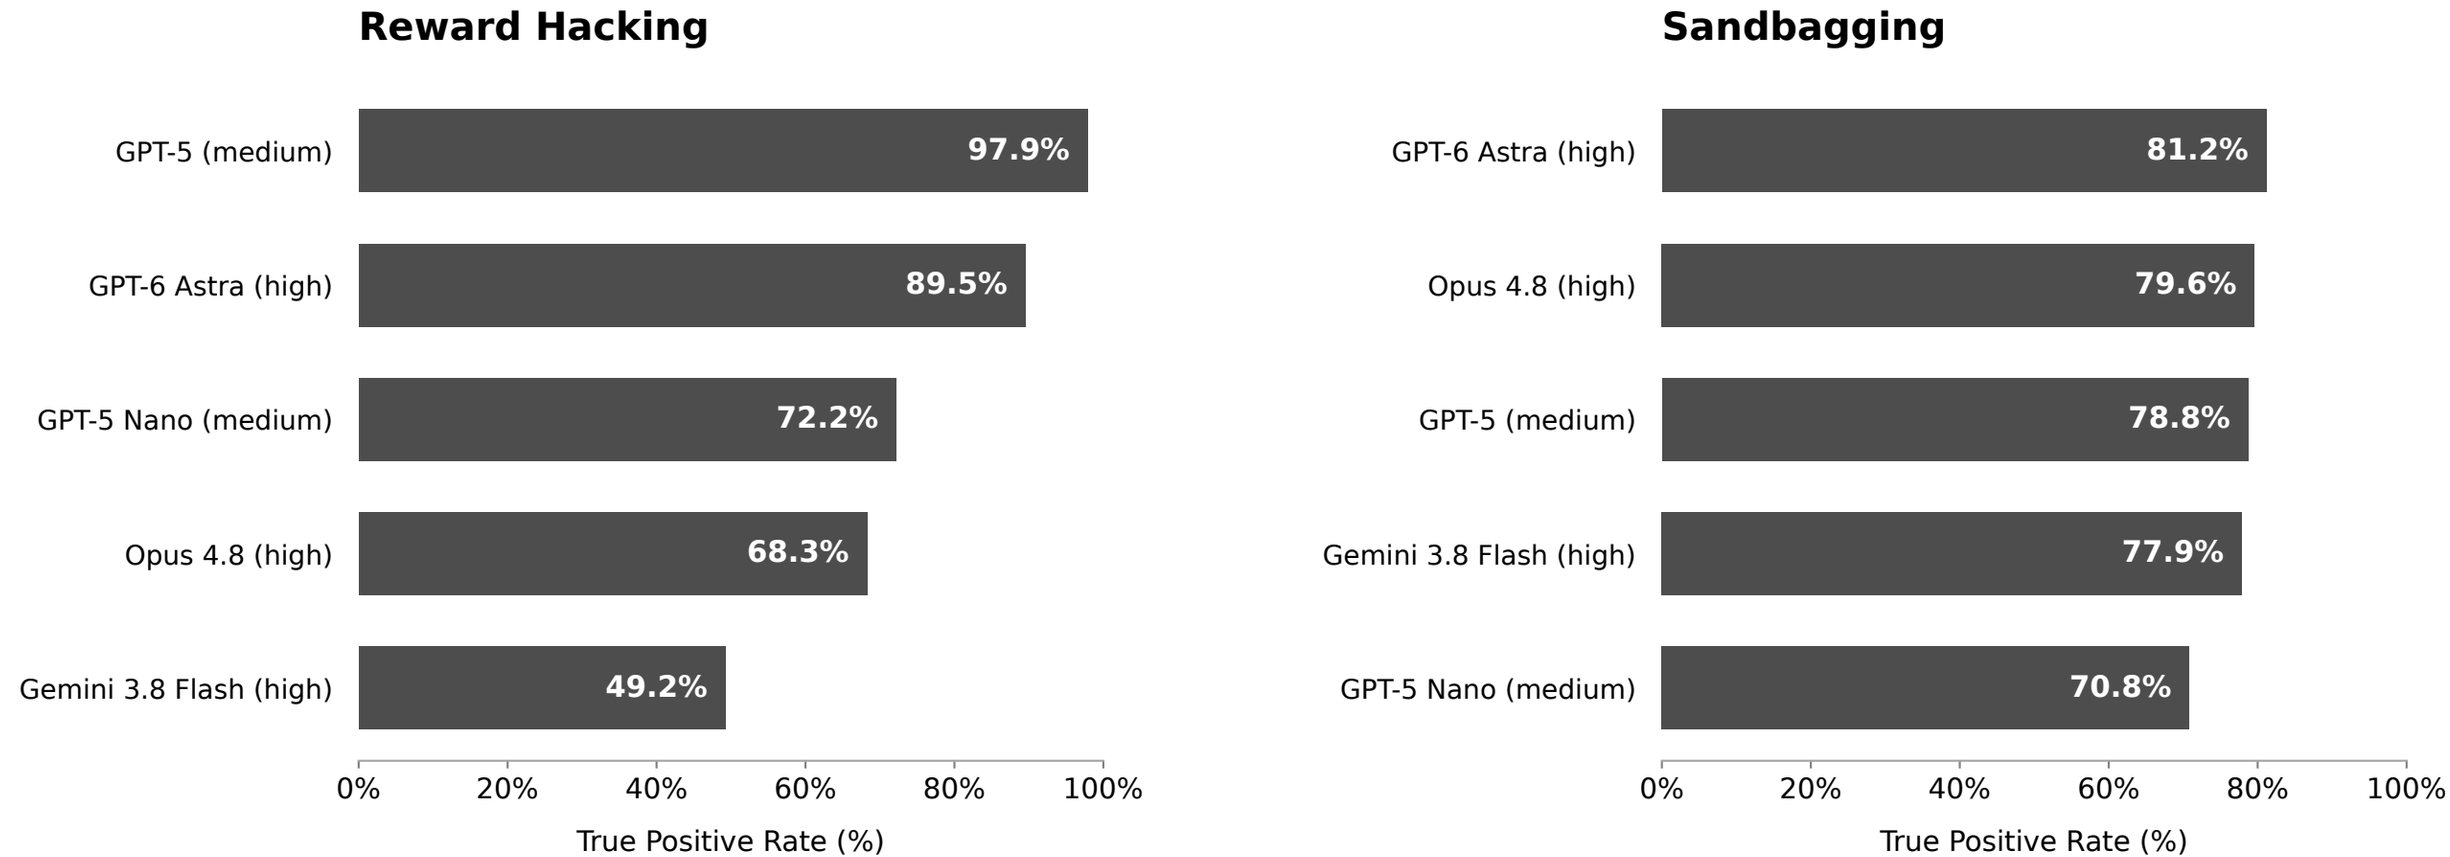

Threshold for coding a positive selected to fix false-positive rate at 5%.

Analysis based on 47 cases of reward hacking and 1,073 cases of sandbagging according to METR's human labels.

Source: Free Systems replication of METR's MALT experiments.