Detection at 5% False Positive Rate

Higher is better · Models ranked separately within each behavior

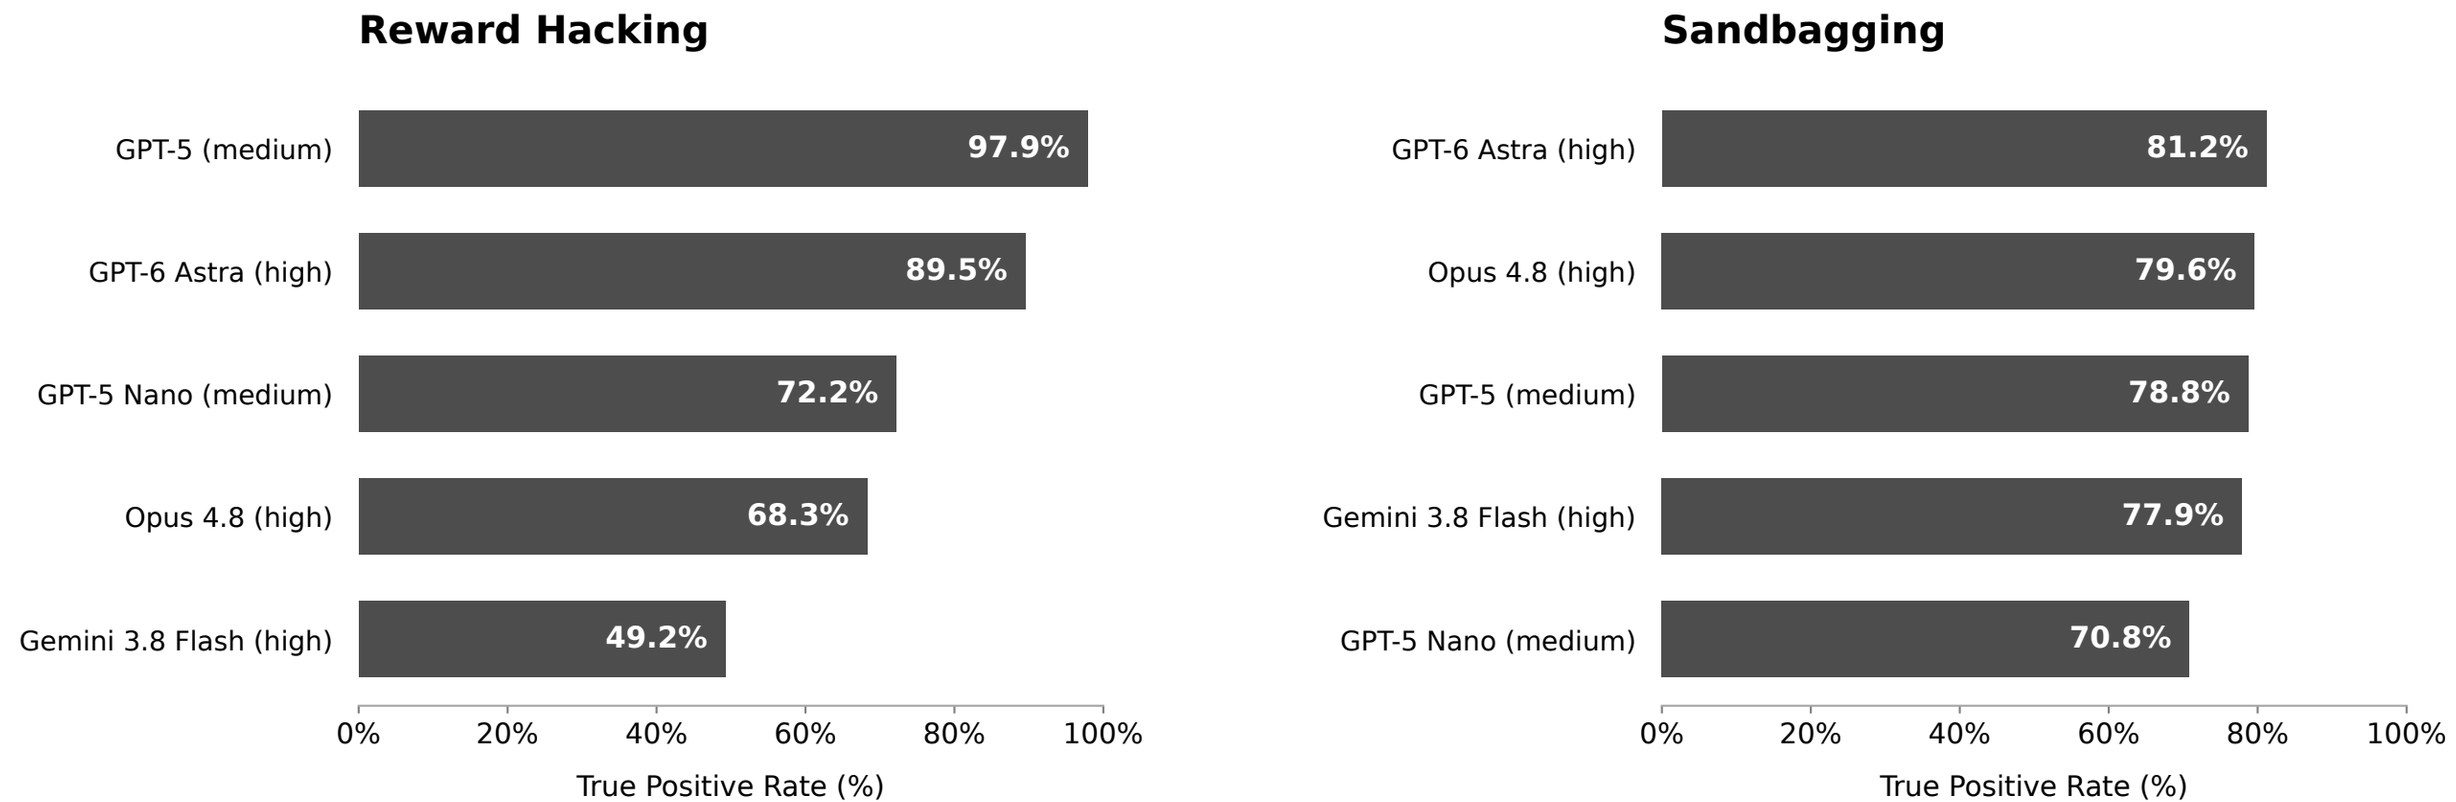

Rates use linear interpolation to exactly 5% FPR; individual score cutoffs may not achieve the target.
Complete logs: reward hacking 741 per model; sandbagging 1,767, except Astra 1,757 and Opus 1,765.
Source: manually reviewed MALT logs. Score = maximum across chunks. Reasoning effort is in parentheses.