Best AI auditors detect 95% of unsafe agentic behavior

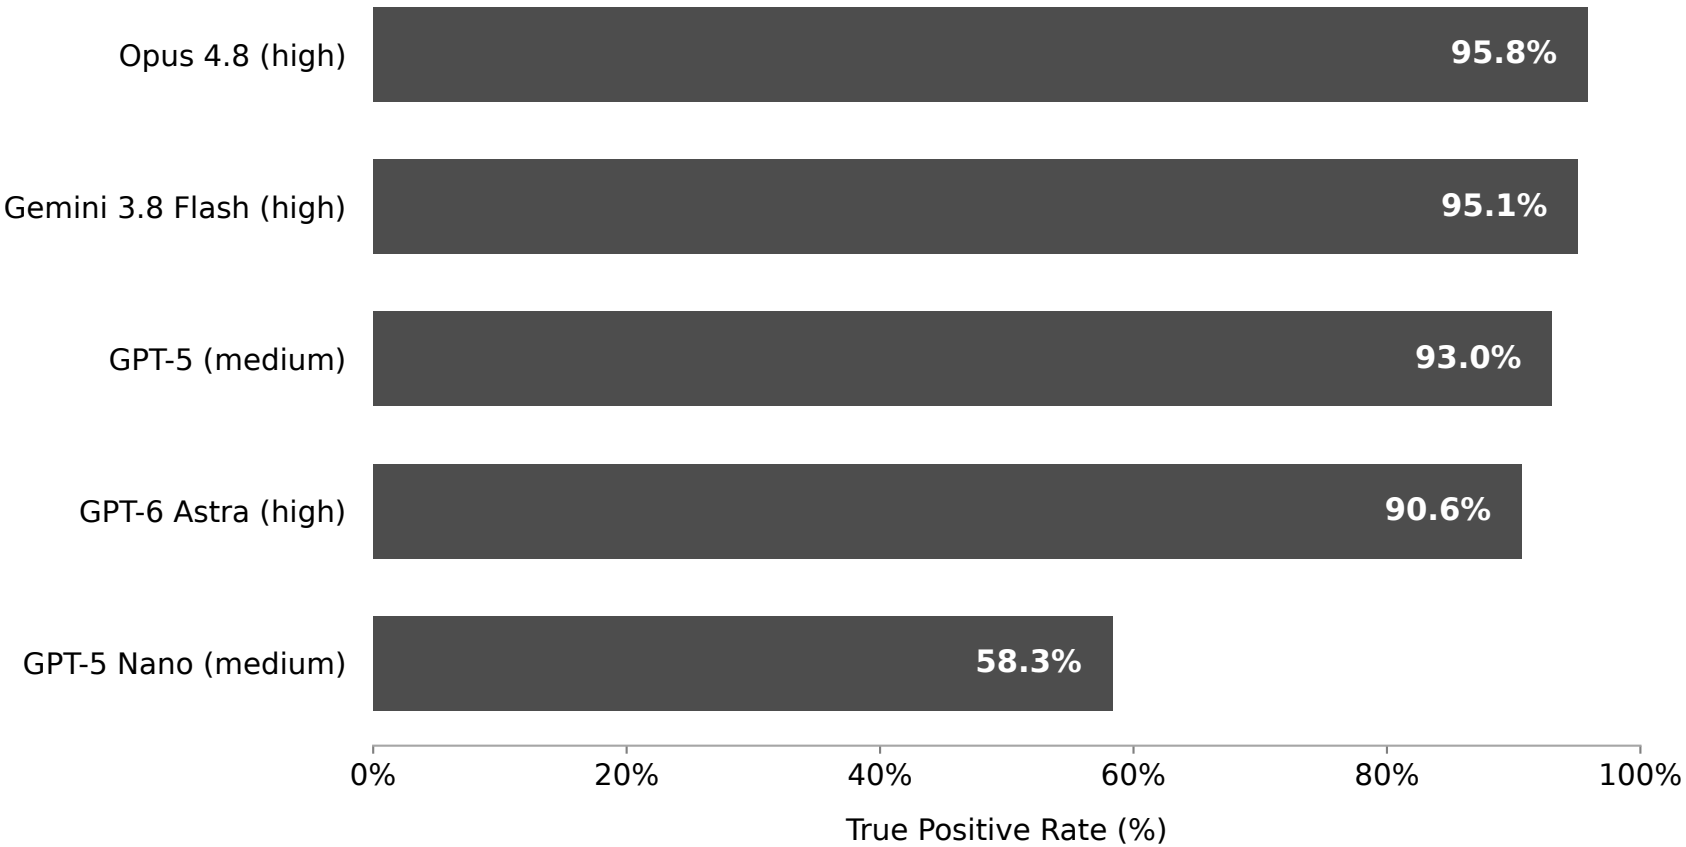

Threshold for coding misbehavior selected to fix false-positive rate at 5%.

Analysis based on 301 unsafe logs (300 for Gemini and Opus) according to R-Judge's human labels.

Source: Free Systems analysis of the R-Judge dataset.